

Discovering, recognizing, and localizing dynamic affordances to adapt to unknown agents

Simon L. Gay
LCIS, University of Grenoble Alpes
Valence, France
simon.gay@lcis.grenoble-inp.fr

Jean-Paul Jamont
LCIS, University of Grenoble Alpes
Valence, France
Jean-Paul.Jamont@lcis.grenoble-inp.fr

Olivier L. Georgeon
EA1598, Lyon Catholic University
LIRIS, Claude Bernard University
Lyon, France
ogergeon@univ-catholyon.fr

Abstract—We present an architecture for self-motivated and environmentally-agnostic agents to integrate non-predictible elements in their emerging model of their environment. This approach aims designing agents that construct their own model of their environment through experience, rather than exploiting pre-coded knowledge. Over time, the agent learns the relation between its perception of objects and the interactions that they afford, in the form of data structures, called signatures of interaction. These signatures are used to recognize and localize objects (or affordances) in surrounding environment, that are stored and tracked in a spatial memory, and used to generate behaviors satisfying agent’s motivational principles. A long term objective of this approach is to study the interaction between multiple agents and the emergence of collaborative or competitive behaviors. Emergence of such behaviors requires the possibility to integrate other agents in the environment model and predict possible actions of these agents. In this paper, we propose a novel architecture of signature of interaction that can define, recognize and localize non-predictible affordances as a first step toward agent intersubjectivity.

Index Terms—developmental learning, interactionism, affordance, autonomous mental development, spatial awareness.

I. INTRODUCTION

We address the problem of how an artificial agent that learns an emergent model of its environment through interaction can acquire knowledge about mobile entities that move freely in the environment (e.g., other agents). We build upon our previous work that learned a model of fixed entities and their geometrical properties. The addition into the environment of entities that move according to their own rules makes the learning even more challenging.

This study is situated within the framework of artificial constructivist learning [1] and enactive learning [2]. In this framework, the learning occurs through the enaction of control loops that implement Piagetian *sensorimotor schemes* [3]. We initialize the agent with a predefined set of control loops called *primitive interactions*. We associate each primitive interaction with a predefined numerical *valence*. We design the agent’s software so that the agent seeks to enact interactions that have positive valences, and to avoid enacting interactions that have negative valences. Such software endows the agent with a kind of self motivation that we have called *interactional motivation* [4]. This view on self motivation relates to the notion of *intrinsic motivation* of artificial agents for *developmental*

learning [5]. For a complete description, we refer the reader to our previous papers in which we have called this approach *Radical Interactionism* (RI) [6] and *artificial interactionism* [7]. Overall, the learning is *unsupervised*. There are no human-defined labels attached with perceptions or with categories of entities, no pre-coded ontology of the world, no semantics associated with action data or sensory data.

In designing agents that autonomously learn to interact with other agents, we pursue the long-term goal of allowing the emergence of collective behaviors within groups of artificial agents. Generating collective behaviors involving other agents requires overcoming two main problems:

- 1) learning to define, recognize and localize other agents that make individual movements in the environment,
- 2) inferring the intentions of these other agents based on their own environmental contexts, which is related to the intersubjectivity problem.

This paper focuses on the resolution of the first problem. It is subdivided as follow: Section II summarizes and formalizes the Radical Interactionism model and the architecture constructing the environment model. Section III presents a model for defining and recognizing probabilistic affordances, and Section IV presents a mechanism to recognize and localize distant probabilistic affordances in space. Finally, Section V encompasses some conclusive remarks and future development of the intersubjectivity problem.

II. THE RADICAL INTERACTIONISM (RI) MODEL

In contrast with most machine learning approaches, an RI agent cannot directly detect the state of its environment through perceptual data. The agent’s input data is not a percept of the state of the environment but an outcome of a control loop. An RI agent learns and exploits regularities in sequences of the control loops offered by its coupling with its environment. The learning mechanism differs from reinforcement learning (e.g., as it is typically implemented in a Partially Observable Markov Decision Process, POMDP) by the fact that RI agents have no reward defined as a function of the system’s state. Our goal is not to design agents that reach predefined goal states or maximize a reward value, but to generate open-ended learning and collective behaviors.

Let I be the set of predefined primitive interactions (pointers to predefined control loops), and $\nu_i \in \mathbb{R}$ the predefined valence

of primitive interaction $i \in I$. At the beginning of step t , the agent selects an *intended interaction* $i_t \in I$. An example consists in moving forward for a predefined amount of time while detecting no obstacle. At the end of step t , the agent is informed of the *enacted interaction* $e_t \in I$ that was actually enacted. If $i_t = e_t$ then the enaction is a success. The agent *did* move forward without detecting an obstacle. Otherwise, the tentative enaction of i_t was a failure. For example, the agent actually enacted another interaction consisting in bumping into an obstacle, which may have a negative valence. An RI agent learns to anticipate the results of its interactions, and tries to enact interactions that have high valences.

To help discovering spatial regularities of the environment, we have proposed the Parallel-RI (PRI) model which allows the simultaneous enaction of multiple interactions. At the end of step t , the agent receives a set $E_t = \{e_{1,t}, \dots, e_{k,t}\} \subset I$ of k enacted interactions instead of a single enacted interaction. For example, when the agent is moving, each enacted interaction can reflect the relative displacement of a salient point in the visual field. This allows the computation of optical flow used to infer the agent's displacements.

Previous PRI experiments showed that the agent was able to discover and recognize static objects through the visual interactions that they afforded (relative displacement of salient points). The agent learned to localize distant affordances (possibilities of enacting an interaction, [9]), and to store and keep track of them in an emergent structure that played the role of spatial memory. The agent generated an implicit context in spatial memory, and used it to generate behaviors in accordance with its interactional motivation. A subsequent model also showed the possibility to identify objects that moved in a straight line, by considering sequences of consecutive interactions [8]. These models, however, worked poorly with entities that made random movements (other agents), because it was not possible to associate the presence of an observed entity with a specific affordance. The present study addresses this limitation.

III. INTEGRATING PROBABILISTIC AFFORDANCES

This section explains the *signature mechanism* by which the agent estimates the possibilities of enacting interactions in a given context. If the experimental settings satisfy the Markov property, we can use the set $E_{t-1} \subset I$ of the interactions enacted at step $t-1$ to define the current context upon which the agent estimates the probability of success of future interactions. Otherwise, we must chain sequences of interactions to extend the temporal dependency.

Formally, a *signature of interaction* is a function $S_i : \mathcal{P}(I) \rightarrow [-1; 1]$, where $\mathcal{P}(I)$ is the partition of I , i.e., the set of all possible contexts. $S_i(E_{t-1}) \in [-1, 1]$ gives the estimated possibility of successfully enacting interaction i at step t : 1 for certainty of success, -1 for certainty of failure. The agent adjusts S_i each time i succeeds or fails to improve the estimation.

We say that a context *affords* (or presents an *affordance* for) interaction i if it makes possible the successful enaction

of i . This terminology complies with Gibson's, and with further formalizations by Chemero [10] and Stoffregen [?]. A signature must be *reversible* : it must be possible to define a function $\hat{S}_i : \{1; -1\} \rightarrow \mathcal{P}(\mathcal{P}(I))$ such that $\hat{S}_i(1)$ gives a set of contexts affording i , and $\hat{S}_i(-1)$ gives the set of contexts preventing the enaction of i . These sets of context should be *minimal* (i.e. $\nexists E_1, E_2 \in \hat{S}_i(x), x \in \{1; -1\} / E_1 \subset E_2$). We note $C_k^{S_i} \in S_i$ a minimal context affording i .

However, previous implementations fails to integrate probabilistic affordances, as the presence of the affordance may still lead to a failure of the interaction. We thus propose to first separate certainty of presence of affordance and probability of success.

A. Separating certainty and probability

When the agent tests an interaction, three cases can be defined from an external observer's perspective:

- 1) The affordance is present at the right place, and the agent can enact the interaction successfully. (e.g. a prey is in front of the agent, and the agent could catch the prey).
- 2) The affordance is present, but the interaction fails. (e.g. a prey is in front of the agent, but the prey moved and the agent missed catching the prey).
- 3) The affordance is absent, leading to a failure of the interaction. (e.g. there is no prey in front of the agent, leading to a certain failure of catching a prey).

From the agent's perspective, the results of situations 2 and 3 cannot be distinguished, as they have the same result. Moreover, situations 2 prevents the construction of the signature. However, our preliminary tests showed an interesting results: despite remaining negative due to situations 2, the signature predictions are slightly higher in case of situations 1-2 than in situations 3. Indeed, the existence of situations 1 influence the signature, allowing the context of interactions corresponding to the affordance to emerge while remaining insufficient to predict with a positive value.

From this observation, we propose to use the average prediction in case of failure \bar{S}_i^f as a threshold to distinguish situations 2 and 3 after a failure. Then, when the interaction failed in an assumed situation 2 (i.e. $S_i(E_t) > \bar{S}_i^f$), the signature is not reinforced, which limit the influence of situations 2 in signature construction. This principle is applied when \bar{S}_i^f is lower than a threshold, assuming that the signature started to emerge (we used a threshold of -0.9).

Several interactions can be prevented by the presence of an element (e.g. *moving a step forward* is expected to success, until an obstacle appears). In such cases, situations 1' are related to the failure of the interaction, and situations 2' and 3' to its success. As situations 2' and 3' are expected to be more frequent than situations 1', the average of predictions will converge to a positive value. We thus used the average of positive predictions \bar{S}_i^s as a threshold to prevent the signature reinforcement in case of success in an assumed situation 2' (i.e. $S_i(E_t) < \bar{S}_i^s$). This principle is applied when \bar{S}_i^s is greater than a threshold (we used a threshold of 0.9).

It is also possible to define the ratio of interaction success when the prediction $S(E_t)$ is greater than the average $\overline{S^f}$ (or interaction failure when $S(E_t)$ is lower than the average $\overline{S^s}$), implying that the agent is in a situation of type 1 or 2 (or 1' or 2'). This ratio thus measures the probability that the presence of the affordance actually affords successfully the interaction.

B. Separating Possible Contexts

The detection of distant affordances (described in Section IV) requires to differentiate the possible contexts affording an interaction, and attribute a probability to each context.

We propose an adaptation of the signature mechanism, based on multiple neurons. Each signature S_i consists of a set of n neurons N_k^i ; the neuron with the strongest output defines the prediction of the signature. In case of success, the neuron with the strongest output is reinforced, while a failure reinforce all neurons. This competition leads to a specialization of each neuron for a specific context, while they are desensitized from other contexts. Thus, with a sufficient number of neurons, a signature can identify contexts affording its interaction independently.

Formally, a neuron k is defined as a set of weights w_l^k (with $Card(\{w_l^k\}) = Card(I)$), with an output defined as:

$$N_k^i(E_t) = f\left(\sum_l E_t[l] \times w_l^k\right), \quad f(x) = \frac{1}{1 + e^{-x}} \quad (1)$$

where the interactional context E_t takes the form of a vector of size $Card(I)$, with $E_t[k] = 1$ when interaction i_k is enacted as a success and $E_t[k] = 0$ otherwise.

Then, the response of the group is defined as the maximum output, and remapped to a range in $[-1; 1]$:

$$N_i(E_t) = \max_k(N_k^i(E_t)) \times 2 - 1 \quad (2)$$

In order to consider interactions that are afforded by the absence of their affordance instead of their presence, we added an output weight W_i defining the output of the signature:

$$S_i(E_t) = N_i(E_t) \times W_i \quad (3)$$

The weight W_i is restrained in the interval $[-1, 1]$, allowing to inverse the result of the prediction, which makes neurons able to integrate contexts preventing the enaction of i .

The learning process uses a classical gradient descent:

- The enaction result is defined as $R_t = 1$ in case of success and $R_t = -1$ in case of failure.

- The weight W_i is updated as follow:

$$W_i \leftarrow W_i + \Delta^i \cdot (\alpha * N_i - \alpha/2), \quad \text{with } \Delta^i = R_t - S_i(E_t) \text{ and } \alpha \text{ the learning rate.}$$

- In case of a success, only the neuron with higher output is reinforced, all neurons are reinforced in case of failure:

$$W_k^i \leftarrow W_k^i + \alpha \cdot \Delta^i \cdot N_k^i, \quad \text{with } \Delta_k^i = (R_t \cdot W_i + 1)/2 - N_k^i$$

To increase the specialization of neurons, in case of a success, weights of the most active neuron connected to failed interaction can be reduced (desensitizing this neuron from other contexts than the current one) and reducing weights of other neurons that are related to successful interactions (desensitizing other neurons from the current context).

A consequence of this implementation is that high weights of neurons characterize contexts affording an interaction. It is also possible to consider weights of a neuron by primary interaction, each group containing a weight related to a primary interaction and weights related to its associated secondary interaction. Thus, a signature S_i can be subdivided into minimal contexts $C_{j,n}^{S_i}$, associated to a primary interaction j and a neuron n . It is then possible to define the ratio of success $p_{j,n}^{s,i}$ and of failure $p_{j,n}^{f,i}$ of each context individually by updating them when primary interaction j is enacted and neuron n has the highest activity, and $S_i(E_t) < \overline{S_i^s}$ or $S_i(E_t) > \overline{S_i^f}$.




Fig. 1. Signature architecture based on multiple neurons. As the signature output relies on the neuron with greatest output, neurons are in competitions, leading to a specialization of each neuron for a specific context. The last function remaps the output to range $[-1; 1]$ and uses a global weight W_i , that can inverse the signature result, allowing the representation of contexts preventing the enaction of an interaction.

C. Implementation

This signature mechanism was tested on an artificial agent moving in a 2-dimensional discrete environment. The sensorimotor possibilities of the agent define a list of five primary interactions, listed below:

- \triangleright move forward by one step,
- \blacktriangleright bump in a solid element, - \curvearrowleft turn left by 90° ,
- \blacktriangleleft eat something edible, - \curvearrowright turn right by 90° ,

Interactions move forward, bump and eat are considered as mutually alternative: intending one of these interactions may lead to the enaction of one of the two others instead.

We add a set of secondary interactions provided by the agent's visual system, that can detect colors among {red, green, blue}, and measure distances. Secondary interactions consist in seeing a red, green or blue element while enacting a primary interaction, at a certain (but initially unknown) position in egocentric space. The interaction bump, that does not produce movement, does not generate visual interactions. We discretize the visual field as a regular grid of 15×8 positions centered on the agents that matches the grid of the environment. We thus define $4 \times 3 \times 15 \times 8 = 1440$ possible secondary interactions. Signatures are implemented as sets of 6 neurons. As we only study the emergence of signatures, we define a unique learning mechanism that foster interactions with low certitude of success or failure (low $|S_i(E_t)|$).

The environment is populated by three types of objects offering spatial regularities that the agent can discover by interacting with them, and characterized by a color that makes them recognizable through agent's sensorimotor system: 1)

wall (green), affording bump, 2) algae (red), that are walk-throughable (and thus useless in the agent's perspective), and 3) fishes (blue), affording eat.




Fig. 2. Test environment. The agent is represented as a grey shark (bottom left), wall as green blocks, algae as red leafs and mobile preys as blue fishes. At each simulation step, the fish has 20% of remaining at current position or to move up, down, left and right. This behavior simulate an agent with unknown decisional mechanism.

The fishes move randomly in the environment: at each simulation step, a fish can move in one of these five directions, with a probability of 20%: immobile, left, right, top, bottom. If the fish cannot move in the selected direction because of an obstacle (wall, alga or other fish), the fish remains at its current position, making the immobile situation slightly more probable than other directions. This random movement simulates agents with unknown behavior.

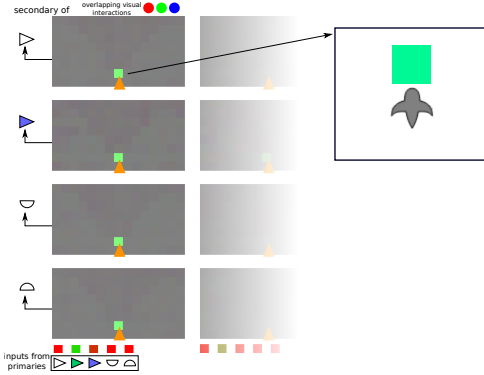


Fig. 3. Signatures of interaction *bump*, recorded after 100 000 simulation steps. A signature is characterized by the weights of 6 formal neurons, each neuron being represented by a column. As the signature identified a unique context, we only represent weights of one neuron. As external observers, we can organize weights of a neuron to make signatures more readable: first, weights associated with primary interactions are represented with five squares below (green for a positive weight, red for a negative weight). Weights associated with secondary interaction are grouped according to their primary interaction, forming the four groups (from top to bottom: forward, eat, turn left, turn right ; bump does not produce visual interactions). Each group is organized to place visual interaction with their associated position in space, relative to the agent (orange triangle). Colors associated with visual interactions are overlapped to generate signatures under the form of a RGB image. Signature of bump identified a context that consist of seeing a green element in front of the agent, which correspond, for an external observer, to the presence of a wall in front of the agent. Bump is also related to the success of bump, as it can bump repeatedly.

We then let the agent behaves in its environment, driven by the signature learning mechanism. Signature of bump emerges

and stabilizes within 5000 simulation steps. The signature is similar to signature obtained in static environments [], and associates the success of bump to the presence of 'seeing a green element in the position in front of the agent', and of a previously enacted 'bump'. The signature thus gathered every interaction allowing to detect the presence of a wall in front of the agent.

Signatures of secondary interactions related to static elements (seeing red and seeing green) progressively stabilize, depending of they frequency of occurrence. After 50 000 simulation steps, most of these signatures stabilized. These signatures are also similar to signatures obtained in static environments []. They designate elements of the same color but on a different position in space. From an external point of view, the spatial offset between the visual interaction and the element designated by its signature matches the movement performed by the enation of its associated primary interaction. This property is used for distant affordance detection [] (details in Section IV).

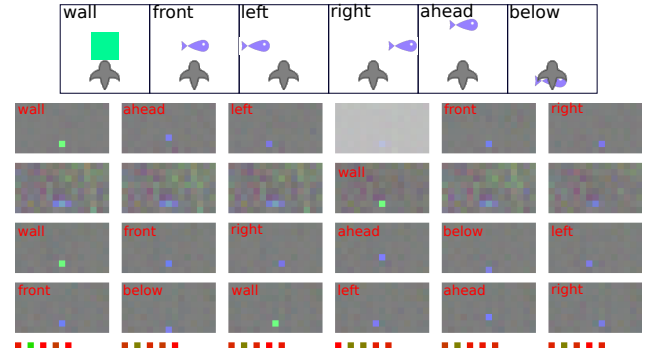


Fig. 4. Signature of *move forward*, recorded after 100 000 simulation steps. Each column represents a neuron of the signature. The weight W is negative: the signature thus represents contexts *preventing* moving forward. The signature identifies six contexts, represented (from an external point of view) above. As a fish cannot be below the agent after forward or eat, only 5 contexts are related to forward primary interaction (greyed context has low weights and is thus unused by the signature). As eat interaction is rarely enacted, contexts related to this primary interaction (second line) are still constructing.

Signatures of interactions related to mobile elements require more steps, as they relate to a larger variety of contexts to identify: at least 45 000 steps are required to identify contexts affording eat and move forward. Signature of interaction eat characterizes five contexts corresponding to the five positions of fish that can lead to a success of eat. Note that the position under the agent does not appear in contexts of interactions associated to move forward, as this situation is not possible. The signature of move forward has a negative weight W . The signature thus shows the affordance that prevents this interaction. The signature designates five contexts associated to the presence of a fish, and one context associated to the presence of a wall in front of the agent.

Signatures of secondary interactions consisting in seeing blue elements designate five contexts, corresponding to the five positions leading to a success of these interactions.




Fig. 5. Signature of eat. The weight W is positive. The signature identifies 5 configurations of fish (front, ahead, left, right, below), 4 in the case of a move forward (as below context cannot be observed). In contexts with fish around front position, we can observe the absence of a green or red element (dark blob) in front of the agent, as this would prevent the prey from moving to this position.




Fig. 6. Signature of secondary interaction seeing a blue element at the position identified with a red square, while moving forward. The weight W is positive. The signature identifies 5 configurations of fish. The signature also indicates the absence of an element that would prevent the fish from moving on the right place, but also elements that could hide the fish (dark blobs).

We also analyze the ratio of successful enaction after a prediction of success. The ratios obtained in contexts implying static objects (such as walls) are close to 1 indicating that the presence of this type of affordance implies the success of the interaction. Ratios obtained with mobile fishes are close to 20%, which correspond to the probability that the fish move in the right direction when the agent tries to eat it. The contexts with a fish in front of the agent is however slightly greater. This can be explained by the fact a prey cannot move a a different position when blocked by a wall or an alga, increasing the probability of eating the fish when in front of the agent. These ratios show that signatures can discover and integrate stochastic properties of the environment.

TABLE I
AVERAGE RATIOS OF SUCCESS WHEN PREDICT A SUCCESS AND FAILURE WHEN PREDICT A FAILURE IN EACH TYPE OF CONTEXT (WALL, FISH IN FRONT, FISH IN OTHER POSITIONS (SURROUNDING))

interaction	ratio of success	ratio of failure
forward (wall)	0.96	0.999
forward (front)	0.25	0.999
forward ('surrounding')	0.19	0.999
bump (wall)	0.97	0.999
eat (front)	0.24	0.999
eat ('surrounding')	0.19	0.999
seeing blue (Fig. 6) (middle)	0.24	0.999
seeing blue (Fig. 6) ('surround.')	0.18	0.999

IV. LOCALIZING DISTANT AFFORDANCES

The detection of distant affordances relies on a property of signatures: a signature of an interaction designates an affordance as sets of interactions $\{j_k\} \in \hat{S}_i(1)$ allowing to detect the presence of this affordance. However, each interaction j_k can have its own signature. It is thus possible to define, from these signatures S_{j_k} , a set of contexts, called *predecessor* $\hat{S}_i^{[j]}$, that, after enacting j , afford i . The *backmove* principle consists in defining the initial predecessor $\hat{S}_i^{\sigma_0} = \hat{S}_i(1)$, with $\sigma_0 = []$. Then, recursively project the predecessor $\hat{S}_i^{\sigma_a}$ with: $\hat{S}_i^{\sigma_{a+1}} = \bigcup_{C_l^i \in \hat{S}_i^{\sigma_a} / j \in C_l^i} \{E \in \mathcal{P}(I) / \forall j_k \in C_l^i, S_{j_k}(E) > 0\}$, with $\sigma_{a+1} = [j, \sigma_a]$ and $C_l^i = \{j_k\}$ contexts of interactions associated to the same primary interaction j . A predecessor \hat{S}_i^σ characterizes a set of contexts that are expected to afford i after enacting sequence σ . Then, when a context $C_l^j \in \hat{S}_i^\sigma$ is observed in E_t , a distant affordance of i is assumed to be present at a position characterized by σ , in egocentric reference.

A. Backmoving a probabilistic affordance

Applying the backmove principle to a signature of a probabilistic affordances would generate a set of predecessor covering all possible future positions of this affordance after enacting a sequence σ , leading to a detection of an affordance through multiple positions in space, which cannot be exploited by a Space Memory. We thus propose to only consider most probable predicted position of an affordance after enacting a sequence σ .

The proposed backmove method introduces a new structure, called projection sequence. The idea is to split predecessors into individual interactions: for each backmove, each sequence considers a unique interaction of S_i^σ . A projection sequence ξ is a structure characterized by:

- a sequence σ of primary interactions, characterizing the movement required to reach the affordance,
- a sequence λ of primary or secondary interactions, that characterize the successive projections from an interaction to an interaction of its signature (principle of backmove).
- a probability p characterizing the probability of enacting i from the partial affordance characterized by the sequence.

The set of projection sequence is constructed as follow: from a signature S_i , a first set of sequence $([], [j_k], p_0)$ is generated for each interaction j_k designated by S_i , where $\sigma = []$ (empty backmove sequence), $\lambda = [j_k]$, and p is the success ratio of the context C_l^i containing j_k . Note that this set characterizes S_i under the form of projection sequences.

Then, the set of sequences is recursively backmoved. A sequence $(\sigma, \lambda, \eta, p)$ leads to interaction $\lambda[0]$. This sequence is backmoved by primitive interaction j associated to $\lambda[0]$ (or by $\lambda[0]$ if primary): from signature $S_{\lambda[0]}$, a set of sequences $([j, \sigma], [j_k, \lambda], [n, \eta], p * p_{S_{\lambda[0]}^{j,n}})$ is generated, for each interaction j_k designated by $S_{\lambda[0]}$.

A sequence ξ_1 is removed from the list if it exists another sequence ξ_2 with $p_{\xi_2} > p_{\xi_1}$ that have similar properties:

- same backmove sequence ($\sigma_{\xi_1} = \sigma_{\xi_2}$)

- same final interaction ($\lambda_{\xi_1}[0] = \lambda_{\xi_2}[0]$)
- divergence comes from different contexts (i.e. $\exists k/\lambda_{\xi_1}[k] \in C_l^i, \lambda_{\xi_2}[k] \in C_m^i, C_l^i \neq C_m^i$), implying that σ_{ξ_1} and σ_{ξ_2} are related to two exclusive future position of the affordance.

The set of projection sequences of a signature S_i provides, for each interaction $i \in I$, a set of most probable sequence of interactions linking interactions from a context E_t and elements designated by S_i . It is then possible to gather sequences ξ_m with the same σ and $\lambda_m[0] \in E_t$, and reconstruct contexts $\{\lambda_m[k]\}$ for each step k of σ , allowing to predict the most probable evolution of an affordance position.

B. Detection of distant affordances

The position of an affordance in space is defined as sequences σ whose enaction would allow to enact the afforded interaction. The sequence is considered regardless of its enactability: the position of an affordance is considered independently from its accessibility.

Projection sequences can be used to detect potential affordances: if a sequence of a signature S_i has its last interaction enacted (i.e. $\lambda[0] \in E_t$), then a part of the context affording i is present at position σ . However, a sequence only characterizes a part of the affordance ; a larger part of the context must be evaluated to confirm the presence of the whole affordance.

The detection of distant affordances of an interaction i starts by selecting projection sequences ξ_k^i of signature S_i whose last element is enacted, defining *candidate* affordances of i . Each candidate ξ_k^i gathers a set of sequences $\Theta_{\xi_k^i} = \{\xi_l^i / \sigma_l = \sigma_k \wedge \lambda_l[0] \in E_t\}$, sharing the same sequence σ and whose last element $\lambda[0]$ is enacted.

The set $E_k^0 = \{\lambda_l[0] / \xi_l \in \Theta_{\xi_k^i}\}$ of last interaction of sequences of an ensemble $\Theta_{\xi_k^i}$ represents an ensemble $E_k^0 \subset E_t$ gathering interactions that can intervene in the presence of the affordance. From the context E_k^0 of a candidate ξ_k^i , the following recursive procedure is applied: elements $\lambda[1]$ of sequences are gathered into a candidate context C^1 (duplicate interactions are removed). Then, each element of C^1 is evaluated with its signature on context E^0 . Interactions predicted as a failure are removed from C^1 , and their projection sequences, removed from $\Theta_{\xi_k^i}$. Remaining interactions of C^1 define context E^1 . The process is repeated with C^{l+1} , until sequence σ is completed (or until $\Theta_{\xi_k^i}$ is empty). Interaction i is then predicted using $S_i(E^l)$. If the signature predicts a success, the affordance of i is confirmed at position σ .

Confirmed affordances of the same interaction can then be gathered if they are characterized by the same elements of the context. the different sequences σ thus indicate the different way to reach this affordance.

C. Test Environments

The affordance detection mechanism was tested with signatures recorded after 150 000 simulation steps in the environment described in Section III-2. The projection mechanism is added to the architecture, and used to generate projection sequences, with a maximum length up to 7 interactions.

As we implemented signature with neurons, we had to adapt the projection sequence construction mechanism. First, we only project interactions designated by a signature with a weight with an absolute value that is greater than a threshold, eliminating non-significant weights. Then, we had to cope with low weights that generate sequences with higher probability than weights with higher (and thus more representative of the affordance). We added a new property to projection sequences, the global weight, characterizing the pertinence of the sequence to represent the affordance. This global weight is computed as follow: first sequences have a global weight defined as $w^{global} = W_S \times w_k$. Then each backmove through a weight w_k of a signature S , the global weight of the new sequence is updated as $w^{global} = w^{global} \times W_S \times w_k$. The filter mechanism then compare values $p \times w^{global}$ instead of p alone, offering a good compromise between probability and pertinence of sequences.

The agent is then presented different environment configurations, and an enaction cycle is performed to let the agent perceive its environment through enacted interactions. The sequences of detected affordances are then analyzed.




Fig. 7. Detection of distant affordances. The agent is presented several elements affording its interactions. These distant affordances are detected and localized through sequences of interactions allowing to reach them. Circles show the position and orientation associated to these sequences (red: not moving forward, blue: eat, green: bump). We can notice that the agent ignore the alga, as it as the same interactional property than an empty space.

Figure N shows the detection in a context containing two wall blocks, two fishes and an alga. Sequences localizing static objects (walls) allows to move toward them. Sequences localizing fishes does not reach the position of the prey, but a position just next to it. Indeed, as the agent can eat a fish on its side with only a slightly lower probability than in front of it, the resulting sequence is a compromise between probability and length of the sequence. Alga is ignored, as it has the same property than empty space. Thus, the affordance detection mechanism can still detect and localize distant affordances under the form of sequences of interactions, which can be stored and exploited by the space memory in a similar way than in a static environment.

V. CONCLUSION AND FUTURE WORK

This work propose a new mechanism to enable an environmentally agnostic agent to consider mobile and non-predictable elements in its emergent model of the environment. Results obtained in a simulated environment showed that the agent can still detect position of distant affordances without the notion of space, and define these position in a similar

way than in static environment, allowing the use of a space memory.

Future work will study how the space memory can be used in such a stochastic environment, and how intrinsically motivated decisional mechanism can integrate probabilities of presence of affordance into consideration.

We will also implement these mechanisms in multi-agent contexts, to study the mutual integration of agents in their own environmental model, and how these models can be exploited for generating behaviors solving collaborative tasks, such as coordinate hunting of large preys. We also plan to study the possibility of predicting other agent's intentions through observation of its own context, as a previous implementation of the space memory demonstrated the possibility for reference change, opening intersubjectivity possibilities between agents.

REFERENCES

- [1] Kristinn R. Thorisson, "A new constructivist ai: From manual methods to self-constructive systems," in *Theoretical Foundations of Artificial General Intelligence*, Pei Wang and Ben Goertzel, Eds., vol. 4, pp. 145–171. Atlantis Press, 2012, Series Title: Atlantis Thinking Machines.
- [2] Tom Froese and Tom Ziemke, "Enactive artificial intelligence: Investigating the systemic organization of life and mind," *Artificial Intelligence*, vol. 173, no. 3–4, pp. 466–500, 2009.
- [3] Jean Piaget, *The construction of reality in the child*, New York: Basic Books, 1954.
- [4] Olivier L. Georgeon, James B. Marshall, and Simon Gay, "Interactional motivation in artificial systems: Between extrinsic and intrinsic motivation," in *2012 IEEE International Conference on Development and Learning and Epigenetic Robotics (ICDL)*, 2012, pp. 1–2, IEEE.
- [5] Pierre-Yves Oudeyer, Frederic Kaplan, and Verena Hafner, "Intrinsic motivation systems for autonomous mental development," *IEEE Trans. Evol. Computat.*, vol. 11, no. 2, pp. 265–286, 2007.
- [6] Olivier Georgeon and David Aha, "The radical interactionism conceptual commitment," *Journal of Artificial General Intelligence*, vol. 4, pp. 31–36, 12 2013.
- [7] Mathieu Guillermin and Olivier L. Georgeon, "Artificial interactionism: Avoiding isolating perception from cognition in ai," *Frontiers in Artificial Intelligence*, vol. 5, 2022.
- [8] Simon L. Gay and Salima Hassas, "Autonomous object modeling based on affordances in a dynamic environment," in *6th Annual International Conference on Biologically Inspired Cognitive Architectures, BICA 2015, Lyon, France, November 6-8, 2015*, Olivier L. Georgeon, Ed. 2015, vol. 71 of *Procedia Computer Science*, pp. 150–156, Elsevier.
- [9] James J Gibson, "The theory of affordances," *Perceiving, Acting, and Knowing: Toward an Ecological Psychology*, Hilldale, USA, vol. 1, no. 2, pp. 67–82, 1977.
- [10] Anthony Chemero, "An outline of a theory of affordances," *Ecological Psychology*, vol. 15, no. 2, pp. 181–195, 2003.
- [11] Stevan Harnad, "The symbol grounding problem," *Physica D: Nonlinear Phenomena*, vol. 42, no. 1-3, pp. 335–346, 1990.